

Stack Overflow is a question and answer site for professional and enthusiast programmers. It's 100% free, no registration required.

Take the 2-minute tour



C++ memory ordering



In some tutorial i saw such spin lock implementation

```
class spin_lock
{
    atomic<unsigned int> m_spin ;

public:
    spin_lock(): m_spin(0) {}
    ~spin_lock() { assert( m_spin.load(memory_order_relaxed) == 0);}

    void lock()
    {
        unsigned int nCur;
        do { nCur = 0; }
        while ( !m_spin.compare_exchange_weak( nCur, 1, memory_order_acquire ));
    }
    void unlock()
    {
        m_spin.store( 0, memory_order_release );
    }
};
```

Do we really need `memory_order_acquire` / `memory_order_release` tags for `compare_exchange_weak` and `store` operations respectively? Or `memory_order_relaxed` is sufficient in this case as there is no Synchronizes-With relation?

Update: Thank you for the explanations! I was considering `spin_lock` without context in which it is used.

c++ c++11 memory-model

edited Dec 24 '13 at 7:12

asked Dec 24 '13 at 6:33



sliser

927 4 9

3 "there is no Synchronizes-With relation" Yes, there is. Eg: Suppose there is a `spin_lock` protect a shared resource, and thread A use `spin_lock.lock()` to acquire the resource, while the thread B is unlocking the lock. There need a memory order to prevent instruction reorder. – KaiWen Dec 24 '13 at 6:41

I would implement spinlocks with `atomic_flag`, since it's the only guaranteed portable lock-free atomic type. – Casey Dec 24 '13 at 14:10

2 Answers

With regard to `memory_order_acquire` / `memory_order_release`: consider the shared data (a.k.a. memory locations, or resource) that you are using the lock to protect. I'll call that "the protected data".

The code needs to ensure that when `lock()` returns, any access to the protected data by the caller will read valid values (i.e. not stale values). `memory_order_acquire` ensures that the appropriate acquire memory barrier is inserted so that subsequent reads of the protected data (via the local CPU cache) will be valid. Similarly, when `unlock()` is called, `memory_order_release` is needed to ensure that the appropriate memory barrier is inserted so that other caches are correctly synchronised.

Some processors don't require acquire/release barriers and might, for a theoretical example, only require a full barrier in `lock()`. But the C++ concurrency model needs to be flexible enough to support many different processor architectures.

`memory_order_relaxed` is used in the destructor because this is just a sanity check to ensure that the lock isn't currently held. Destructing the lock does not confer any synchronisation

semantics to the caller.

answered Dec 24 '13 at 7:02



Ross Bencina
919 2 12

{ USE STACK OVERFLOW TO
FIND THE BEST DEVELOPERS }

 stackoverflowcareers

Yes `memory_order_acquire` / `memory_order_release` is needed.

You use locks to protect shared data between multiple threads. if you don't use `memory_order_release` in `unlock` method, any write on shared data can be reordered after `unlock` method. Also if you don't use `memory_order_acquire` in `lock` method any read on shared data can be reordered before `lock` method. so you need `acquire/release` to protect shared data between threads.

```
spinLock.lock() // use acquire here, so any read can't reordered before `Lock`
```

```
// Writes to shared data
```

```
spinLock.unlock() // use release here, so any write can't reordered after `unlock`
```

with `acquire/release` all writes to shared data will be visible to threads that lock `spinlock`.

edited Dec 24 '13 at 7:06

answered Dec 24 '13 at 7:01



MohammadRB
1,083 1 4 15